

Raw-return MLP: validation objective vs. training transitions

