

Training-time optimisation diagnostics: Raw-return LSTM

Per-PPO-update diagnostics from the rollout/optimisation step, NOT the deterministic validation objective -- these measure whether/how the policy is still learning, not how well it performs.

