



Training-time learning curve: per-update rollout return_mean vs. transitions

GAE-bootstrapped per-timestep return from PPO's rollout buffer, NOT the deterministic validation objective -- separate y-axis per architecture, scale not comparable across panels or to the analytical/clone benchmarks

Hamilton (diagnostics-only retrain -- not the official training run, see docstring)



Raw-return MLP



Raw-return LSTM

